

Modification of the Database

- The content of the database may be modified using the following operations:
 - Deletion
 - Insertion
 - Updating
- All these operations are expressed using the assignment operator.
- **Deletion**
 - A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
 - Can delete only whole tuples; cannot delete values on only particular attributes. A deletion is expressed in relational algebra by:

$$r \leftarrow r - E$$

where r is a relation and E is a relational algebra query.

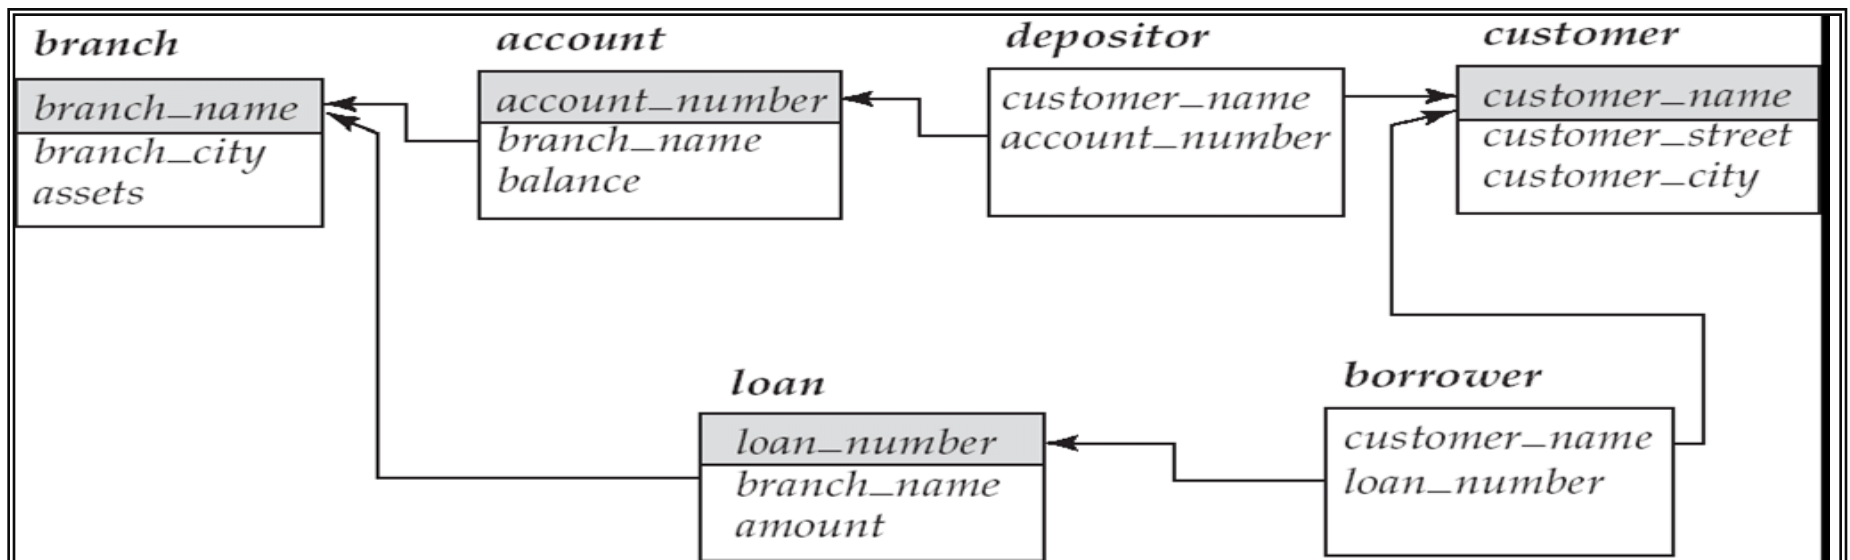
Deletion ...contd.

- Delete all account records in the Perryridge branch.

$account \leftarrow account - \sigma_{branch_name = "Perryridge"}(account)$

- Delete all loan records with amount in the range of 0 to 50

- Delete all accounts at branches located in Needham.



Deletion ...contd.

- Delete all account records in the Perryridge branch.

$account \leftarrow account - \sigma_{branch_name = "Perryridge"}(account)$

- Delete all loan records with amount in the range of 0 to 50

$loan \leftarrow loan - \sigma_{amount \geq 0 \text{ and } amount \leq 50}(loan)$

- Delete all accounts at branches located in Needham.

$r_1 \leftarrow \sigma_{branch_city = "Needham"}(account \bowtie branch)$

$r_2 \leftarrow \Pi_{account_number, branch_name, balance}(r_1)$

$r_3 \leftarrow \Pi_{customer_name, account_number}(r_2 \bowtie depositor)$

$account \leftarrow account - r_2$

$depositor \leftarrow depositor - r_3$

Insertion

▪ To insert data into a relation, we either:

- specify a tuple to be inserted

- write a query whose result is a set of tuples to be inserted in relational algebra, an insertion is expressed by:

$$r \leftarrow r \cup E$$

where r is a relation and E is a relational algebra expression.

- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.

- Example:

- Insert information in the database specifying that Smith has Rs1200 in account A-973 at the Perryridge branch.

- Provide as a gift for all loan customers in the Perryridge branch, a Rs200 savings account. Let the loan number serve as the account number for the new savings account.

- Insert information in the database specifying that Smith has Rs1200 in account A-973 at the Perryridge branch.

$account \leftarrow account \cup \{("A-973", "Perryridge", 1200)\}$

$depositor \leftarrow depositor \cup \{("Smith", "A-973")\}$

- Provide as a gift for all loan customers in the Perryridge branch, a Rs200 savings account. Let the loan number serve as the account number for the new savings account.

$r_1 \leftarrow (\sigma_{branch_name = "Perryridge"}(borrower \quad loan))$

$account \leftarrow account \cup \Pi_{loan_number, branch_name, 200}(r_1)$

$depositor \leftarrow depositor \cup \Pi_{customer_name, loan_number}(r_1)$

Updating

- A mechanism to change a value in a tuple without changing *all* values in the tuple
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_l}(r)$$

- Each F_i is either
 - The I^{th} attribute of r , if the I^{th} attribute is not updated, or,
 - If the attribute is to be updated F_i is an expression, involving only constants and the attributes of r , which gives the new value for the attribute

Update...contd.

- Make interest payments by increasing all balances by 5 percent.

account $\leftarrow \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.05}(\text{account})$

- Pay all accounts with balances over Rs 10,000, 6 percent interest and
- pay all others 5 percent

account $\leftarrow \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.06}(\sigma_{BAL > 10000}(\text{account}))$
 $\cup \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.05}(\sigma_{BAL \leq 10000}(\text{account}))$

Structured Query Language (SQL)

History

- **IBM SEQUEL language developed as part of System R project at the IBM San Jose Research Laboratory**
- **Renamed Structured Query Language (SQL)**
- **ANSI and ISO standard SQL:**
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant)
 - 8.0 for SQL Server 2000.
 - SQL:2003
 - 9.0 for SQL Server 2005.
 - 10.0 for SQL Server 2008.
 - 10.5 for SQL Server 2008 R2.
 - 11.0 for SQL Server 2012.
 - 12.0 for SQL Server 2014
 - 13.0 for SQL Server 2016
 - 14 for SQL Server 2017
 - 15 for SQL Server 2019
 - 16 for SQL Server 2022
- **Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.**

Data Definition Language

- **Allows the specification of:**
 - The schema for each relation including attribute types
 - Integrity constraints
 - Authorization information for each relation.
- **Non-standard SQL extensions also allow specification of:**
 - The set of indices to be maintained for each relations.
 - The physical storage structure of each relation on disk.

Create Table Construct

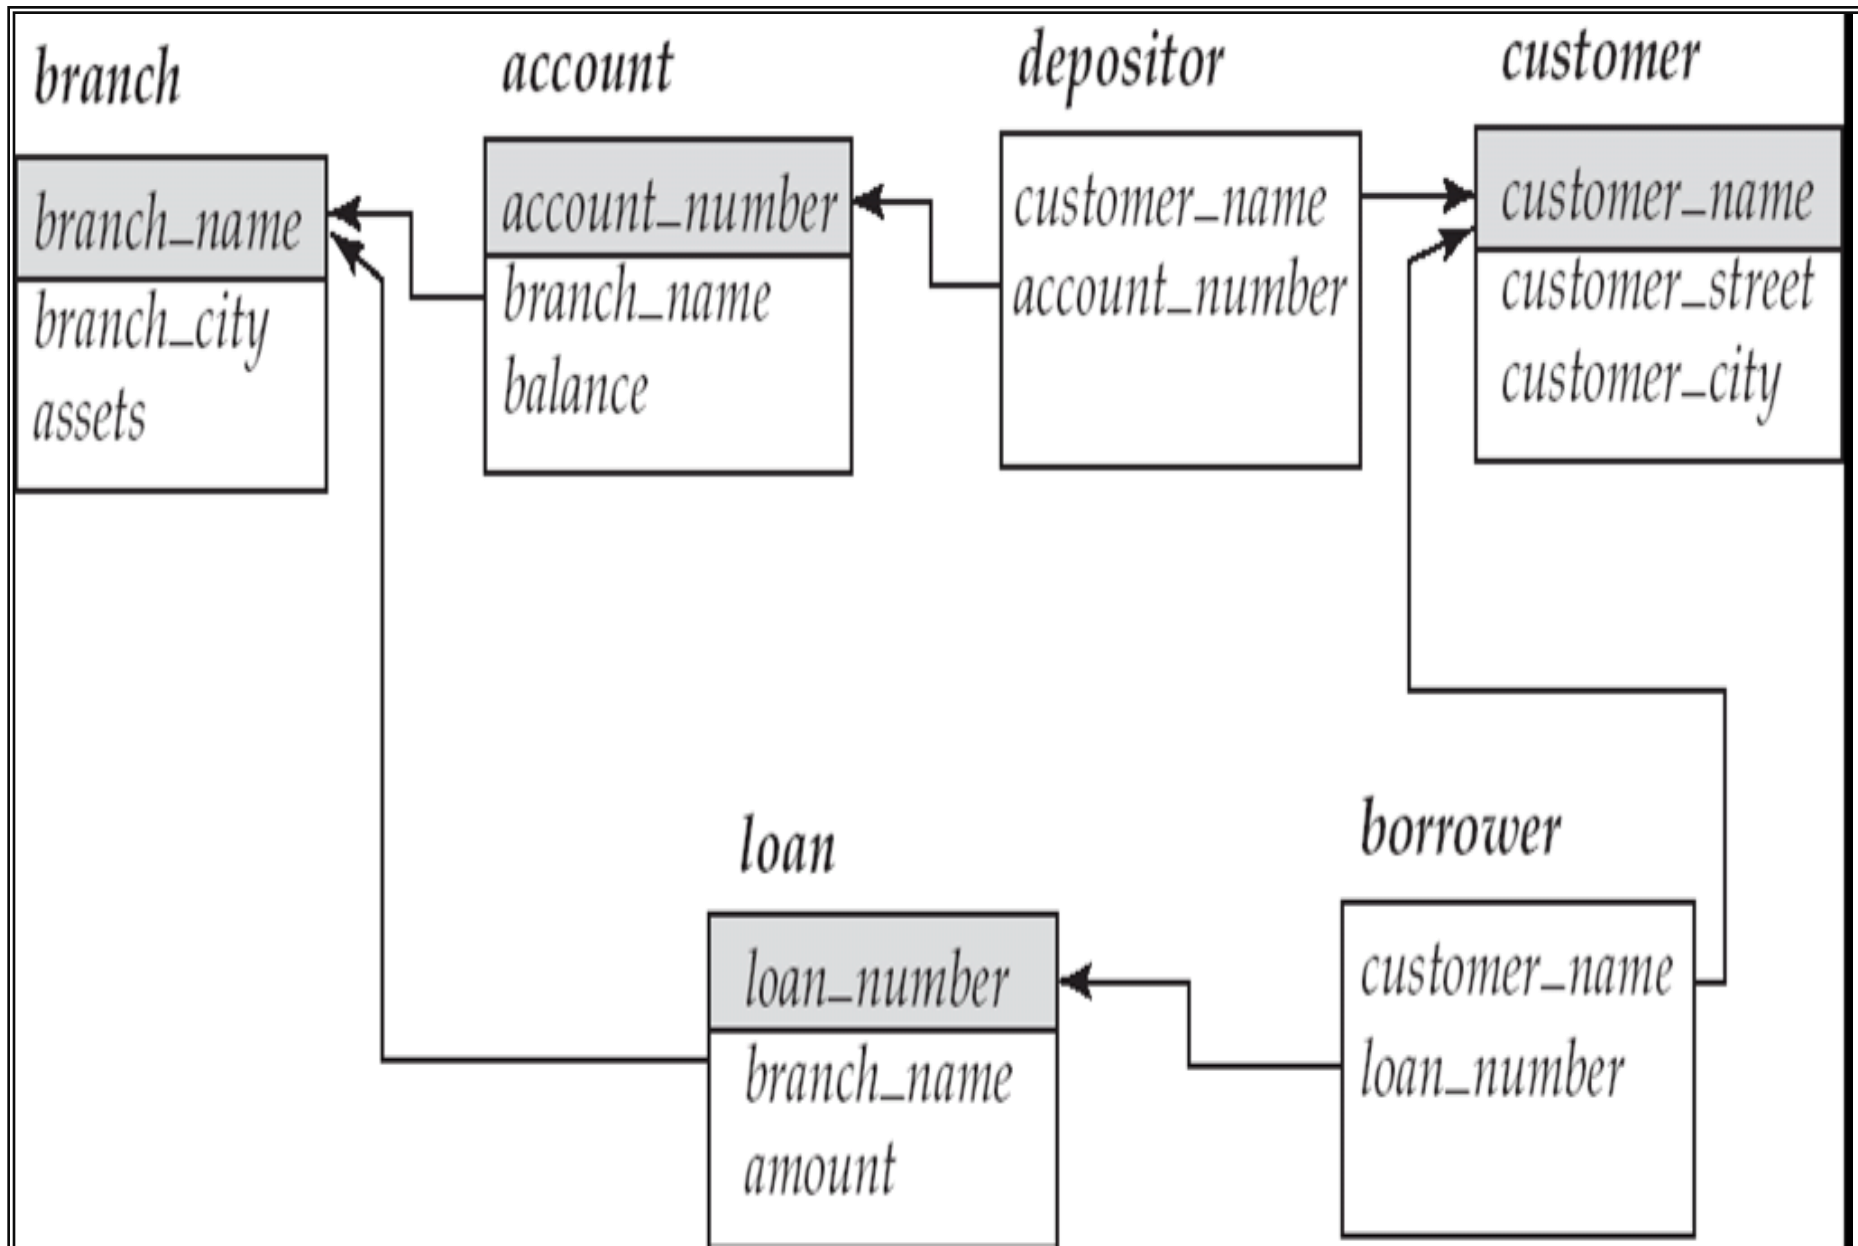
- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,  
                (integrity-constraint1),  
                ...,  
                (integrity-constraintk))
```

- *r* is the name of the relation
- each *A_i* is an attribute name in the schema of relation *r*
- *D_i* is the data type of attribute *A_i*

- Example:

```
create table branch  
    (branch_name char(15),  
    branch_city  char(30),  
    assets        integer)
```



Referential Integrity in SQL

```
create table account  
(account_number    char(10),  
branch_name    char(15),  
balance        integer,  
primary key (account_number),  
foreign key (branch_name) references branch )
```

```
create table depositor  
(customer_name    char(20),  
account_number    char(10),  
primary key (customer_name, account_number),  
foreign key (account_number ) references account,  
foreign key (customer_name ) references customer )
```

```
create table customer  
(customer_name    char(20),  
customer_street    char(30),  
customer_city    char(30),  
primary key (customer_name ))
```

```
create table branch  
(branch_name    char(15),  
branch_city    char(30),  
assets        numeric(12,2),  
primary key (branch_name ))
```

Domain Types in SQL

- **char(*n*)**: Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**: Variable length character strings, with user-specified maximum length *n*.
- **Int**: Integer (a finite subset of the integers that is machine-dependent).
- **Smallint**: Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**: Fixed point number, with user-specified precision of *p* digits, with *n* digits to the right of decimal point.
- **real, double precision**: Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**: Floating point number, with user-specified precision of at least *n* digits.

Integrity Constraints on Tables

- **not null**
- **primary key** (A_1, \dots, A_n)
- Example: Declare *branch name* as the primary key for *branch*

```
create table branch
(branch_name      char(15),
 branch_city char(30) not null,
 assets          integer,
primary key (branch_name))
```

primary key declaration on an attribute automatically ensures **not null** in SQL-92 onwards, needs to be explicitly stated in SQL-89

Basic Insertion and Deletion of Tuples

- Newly created table is empty

- Add a new tuple to *account*

insert into *account*

values ('A-9732', 'Perryridge', 1200)

– Insertion fails if any integrity constraint is violated

- Delete *all* tuples from *account*

delete from *account*

Drop and Alter Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database.

drop table

- The **alter table** command is used to add attributes to an existing relation:

alter table r add A D

where A is the name of the attribute to be added to relation r and D is the domain of A .

- All tuples in the relation are assigned *null* as the value for the new attribute.

- The **alter table** command can also be used to drop attributes of a relation:

alter table r drop A

where A is the name of an attribute of relation r

- Dropping of attributes not supported by many databases

Basic Query Structure

- General form of SQL query :

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

- A_i represents an attribute
- R_i represents a relation
- P is a predicate.

- This query is equivalent to the relational algebra expression.

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

- The result of an SQL query is a relation.

The select Clause

- The **select** clause list the attributes desired in the result of a query
 - corresponds to the projection operation of the relational algebra

- Example: Find the names of all branches in the *loan* relation:

```
select branch_name  
from loan
```

- In the relational algebra, the query would be:

$$\prod_{branch_name} (loan)$$

- To force the elimination of duplicates, insert the keyword **distinct** after select.

- Find the names of all branches in the *loan* relations, and remove duplicates

```
select distinct branch_name  
from loan
```

- The keyword **all** specifies that duplicates not be removed.

```
select all branch_name  
from loan
```

The select ClauseContd.

- An asterisk in the select clause denotes “all attributes”

```
select *  
from loan
```

- The **select** clause can contain arithmetic expressions involving the operation, +, −, *, and /, and operating on constants or attributes of tuples.

- E.g.:

```
select loan_number, branch_name, amount *100  
from loan
```

The where Clause

- The **where** clause specifies conditions that the result must satisfy
 - Corresponds to the selection predicate of the relational algebra.
- To find all loan number for loans made at the Perryridge branch with loan amounts greater than Rs1200.

select *loan_number*
from *loan*
where *branch_name* = 'Perryridge' and *amount* > 1200
- Comparison results can be combined using the logical connectives and, or, and not.

The from Clause

- The **from clause lists** the **relations** involved in the query
 - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product of *borrower X loan*
select *
from *borrower, loan*
- **Find the name, loan number and loan amount of all customers having a loan at the Perryridge branch.**

```
select customer_name, borrower.loan_number, amount  
from borrower, loan  
where borrower.loan_number = loan.loan_number and branch_name = 'Perryridge'
```

Rename Operator

SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

Example: Find the name, loan number and loan amount of all customers; rename the column name *loan_number* as *loan_id*.

```
select customer_name, borrower.loan_number as loan_id, amount
      from borrower, loan
     where borrower.loan_number = loan.loan_number
```

Tuple Variables

- Tuple variables are defined in the from clause via the use of the as clause
- Find the customer names and their loan numbers and amount for all customers having a loan at some branch.

```
select customer_name, T.loan_number, S.amount  
      from borrower as T, loan as S  
     where T.loan_number = S.loan_number
```

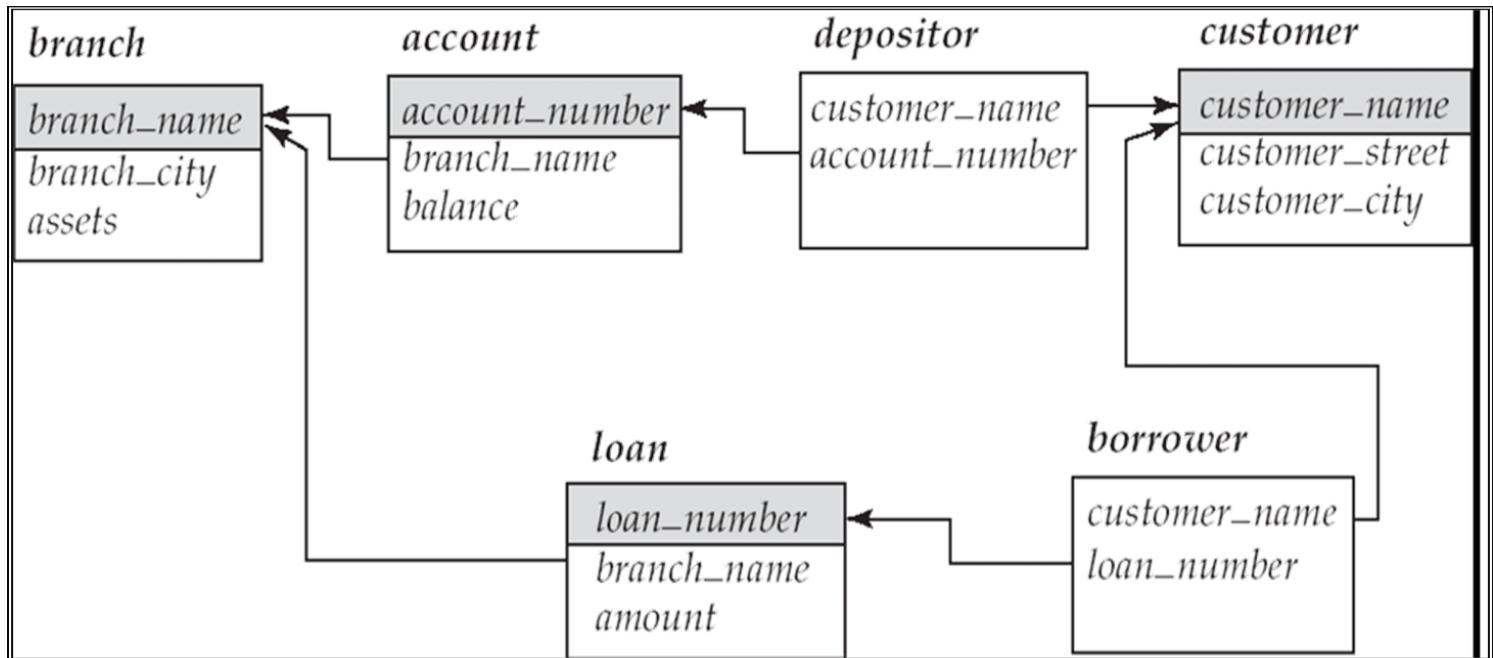
Find the names of all branches that have greater assets than some branch located in Brooklyn.

```
select distinct T.branch_name  
      from branch as T, branch as S  
     where T.assets > S.assets and S.branch_city = 'Brooklyn'
```

Keyword **as** is optional and may be omitted

borrower as T \equiv *borrower T*

Some database such as Oracle *require* **as** to be omitted



String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator “**like**” uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring.
 - underscore (_). The _ character matches any character.
- Find the names of all customers whose street includes the substring “Main”.

```
select customer_name  
from customer  
where customer_street like '% Main%'
```

Match the name “Main%”

```
like 'Main\%' escape '\'
```

SQL supports a variety of string operations such as

concatenation (using “||”)

converting from upper to lower case (and vice versa)

finding string length, extracting substrings, etc.

String Operations

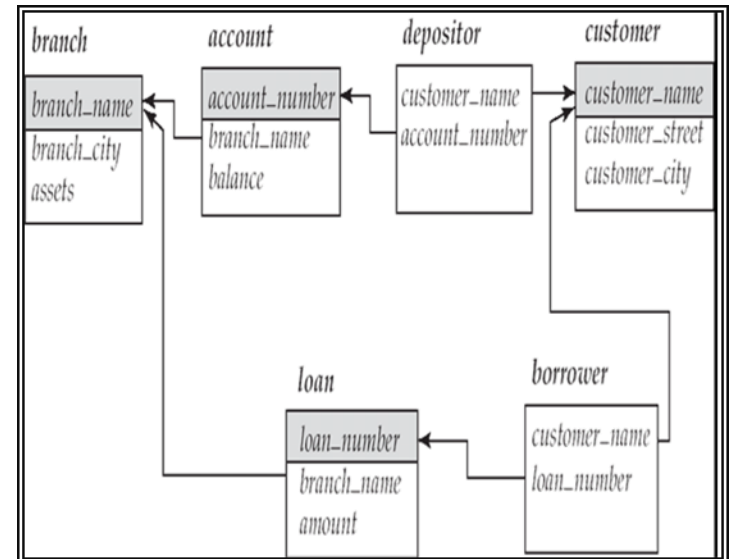
- List in alphabetic order the names of all customers having a loan in Perryridge branch

```
select distinct customer_name
from   borrower, loan
where borrower loan_number = loan.loan_number and
       branch_name = 'Perryridge'
order by customer_name
```

We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

Example:

```
select distinct customer_name
from   borrower, loan
where borrower loan_number = loan.loan_number
       and branch_name = 'Perryridge'
order by customer_name decs
```



Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.
- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2 :

Duplicates ...Contd.

- SQL duplicate semantics:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

Set Operations

- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations
- Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs m times in r and n times in s , then, it occurs:

- $m + n$ times in r **union all** s
- $\min(m, n)$ times in r **intersect all** s
- $\max(0, m - n)$ times in r **except all** s

Set Operations

- Find all customers who have a loan, an account, or both:

(select customer_name from depositor)

union

(select customer_name from borrower)

- Find all customers who have both a loan and an account.

(select customer_name from depositor)

intersect

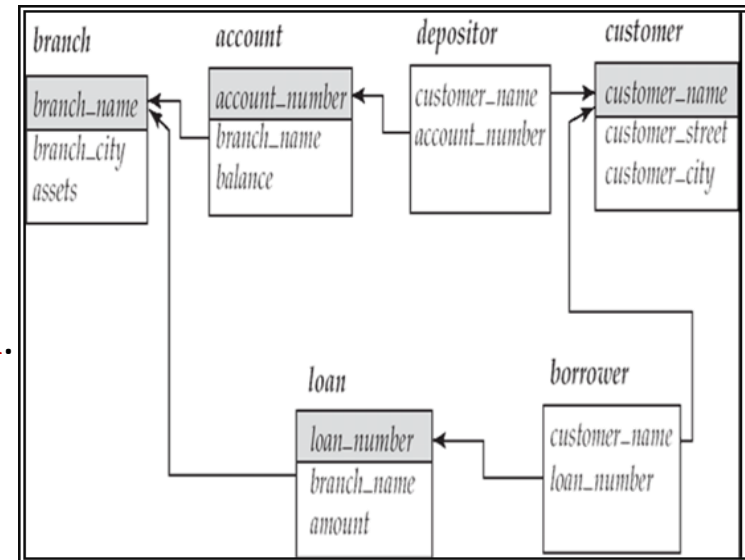
(select customer_name from borrower)

- Find all customers who have an account but no loan.

(select customer_name from depositor)

except

(select customer_name from borrower)



Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value.

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- Find the average account balance at the Perryridge branch.**

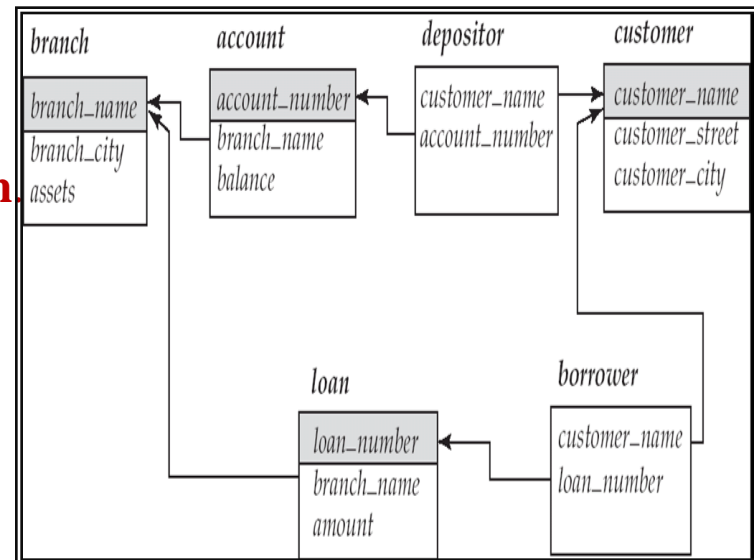
```
select avg (balance)
  from account
 where branch_name = 'Perryridge'
```

- Find the number of tuples in the *customer* relation.**

```
select count (*)
  from customer
```

- Find the number of depositors in the bank.**

```
select count (distinct customer_name)
  from depositor
```



Aggregate Functions – Group By

- Find the number of depositors for each branch.

```
select branch_name, count (distinct customer_name)
  from depositor, account
 where depositor.account_number = account.account_number
 group by branch_name
```

Aggregate Functions – Having Clause

- Find the names of all branches where the average account balance is more than Rs 1,200.

```
select branch_name, avg (balance)
  from account
 group by branch_name
 having avg (balance) > 1200
```

