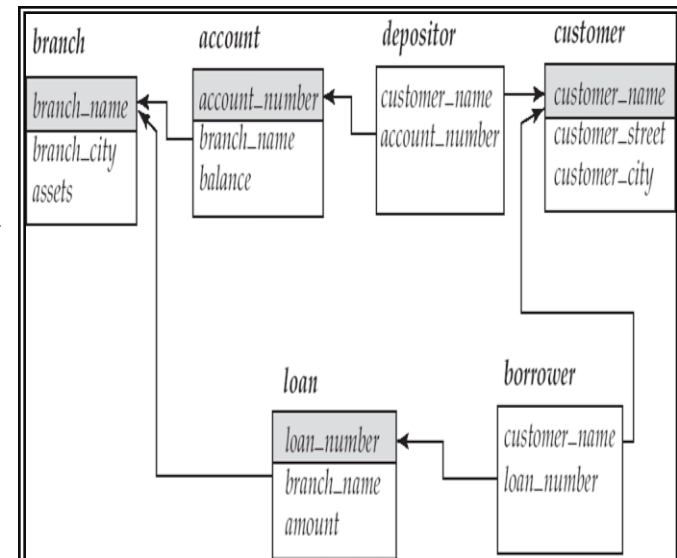


Nested Queries

- SQL provides a mechanism for the nesting of subqueries.
 - A **subquery** is a **select-from-where** expression that is nested within another query.
 - A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.
- Find all customers who have both an account and a loan at the bank.

```
select distinct customer_name
  from borrower
 where customer_name in (select customer_name
                        from depositor )
```



In Construct- Nested Queries

- Find all customers who have a loan at the bank but do not have an account at the bank

```
select distinct customer_name
  from borrower
 where customer_name not in (select customer_name
                             from depositor )
```

- Find all customers who have both an account and a loan at the Perryridge branch

```
select distinct customer_name
  from borrower, loan
 where borrower.loan_number = loan.loan_number and
        branch_name = 'Perryridge' and
        (branch_name, customer_name ) in
        (select branch_name, customer_name
          from depositor, account
         where depositor.account_number =
               account.account_number )
```

“Some” Construct

- Find all branches that have greater assets than some branch located in Brooklyn.

```
select distinct T.branch_name  
  from branch as T, branch as S  
  where T.assets > S.assets and  
        S.branch_city = 'Brooklyn'
```

- Same query using > **some** clause

```
select branch_name  
  from branch  
  where assets > some  
              (select assets  
                from branch  
                where branch_city = 'Brooklyn')
```

“All” Construct

- Find the names of all branches that have greater assets than all branches located in Brooklyn.

```
select branch_name
      from branch
     where assets > all
           (select assets
            from branch
            where branch_city = 'Brooklyn')
```

“Exists” Construct

- Find all customers who have an account at all branches located in Brooklyn.

```
select distinct S.customer_name
from depositor as S
where not exists (
    (select branch_name
     from branch
     where branch_city = 'Brooklyn')
    except
    (select R.branch_name
     from depositor as T, account as R
     where T.account_number = R.account_number and
          S.customer_name = T.customer_name ))
```

- Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- *Note:* Cannot write this query using **= all** and its variants

Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- Find all customers who have at most one account at the Perryridge branch.

```
select T.customer_name
from depositor as T
where unique (
    select R.customer_name
    from account, depositor as R
    where T.customer_name = R.customer_name
    and
    R.account_number = account.account_number and
    account.branch_name = 'Perryridge')
```

- **Find all customers who have at least two accounts at the Perryridge branch.**

```
select distinct T.customer_name
from depositor as T
where not unique (
    select R.customer_name
    from account, depositor as R
    where T.customer_name = R.customer_name and
        R.account_number = account.account_number and
        account.branch_name = 'Perryridge')
```

- Variable from outer level is known as a **correlation variable**

Modification of the Database – Deletion

- **Delete all account tuples at the Perryridge branch**

```
delete from account  
where branch_name = 'Perryridge'
```

- **Delete all accounts at every branch located in the city ‘Needham’.**

```
delete from account  
where branch_name in (select branch_name  
                        from branch  
                        where branch_city = 'Needham')
```


- **Delete the record of all accounts with balances below the average at the bank.**

```
delete from account  
  where balance < (select avg (balance)  
                    from account)
```

- Problem: as we delete tuples from deposit, the average balance changes
- Solution used in SQL:
 1. First, compute **avg** balance and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

Modification of the Database – Insertion

- **Add a new tuple to *account***

```
insert into account  
values ('A-9732', 'Perryridge', 1200)
```

or equivalently

```
insert into account (branch_name, balance, account_number)  
values ('Perryridge', 1200, 'A-9732')
```

- **Add a new tuple to *account* with *balance* set to null**

```
insert into account  
values ('A-777', 'Perryridge', null )
```

Joined Relations – Datasets

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700
<i>loan</i>		

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155
<i>borrower</i>	

Note: borrower information missing for L-260 and loan information missing for L-155

Joined Relations

- **loan inner join borrower on**
loan.loan_number = borrower.loan_number

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>	<i>loan_number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

- **loan left outer join borrower on**
loan.loan_number = borrower.loan_number

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>	<i>loan_number</i>
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	<i>null</i>	<i>null</i>

Joined Relations

- *loan natural inner join borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

loan natural right outer join borrower

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	<i>null</i>	<i>null</i>	Hayes

Find all customers who have either an account or a loan (but not both) at the bank.

```
select customer_name  
  from (depositor natural full outer join borrower )  
 where account_number is null or loan_number is null
```

Derived Relations

- SQL allows a subquery expression to be used in the **from** clause
- **Find the average account balance of those branches where the average account balance is greater than Rs 1200.**

```
select branch_name, avg_balance
from (select branch_name, avg (balance)
      from account
      group by branch_name )
as branch_avg ( branch_name, avg_balance )
where avg_balance > 1200
```

Note : We do not need to use the **having** clause, since we compute the temporary (view) relation **branch_avg** in the **from** clause, and the attributes of **branch_avg** can be used directly in the **where** clause.

View Definition

- A relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a view.
- A view is defined using the create view statement which has the form
create view v as < query expression >
where <query expression> is any legal SQL expression. The view name is represented by v.
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.

- **A view consisting of branches and their customers**

```
create view all_customer as  
    (select branch_name, customer_name  
    from depositor, account  
    where depositor.account_number =  
        account.account_number )  
union  
    (select branch_name, customer_name  
    from borrower, loan  
    where borrower.loan_number = loan.loan_number )
```

■ **Find all customers of the Perryridge branch**

```
select customer_name  
    from all_customer  
    where branch_name = 'Perryridge'
```


Uses of Views

- Hiding some information from some users
 - Consider a user who needs to know a customer's name, loan number and branch name, but has no need to see the loan amount.
 - Define a view

```
(create view cust_loan_data as
select customer_name, borrower.loan_number, branch_name
from borrower, loan
where borrower.loan_number = loan.loan_number )
```
 - Grant the user permission to **read cust_loan_data, but not borrower or loan**
 - Predefined queries to make writing of other queries easier
 - Common example: Aggregate queries used for statistical analysis of data

Processing of Views

- When a view is created
 - Query expression is stored in the database along with the view name
 - Expression is substituted into any query using the view
- Views definitions containing views
 - One view may be used in the expression defining another view
 - A view relation v_1 is said to **depend directly** on a view relation v_2 if v_2 is used in the expression defining v_1
 - A view relation v_1 is said to **depend on** view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2
 - A view relation v is said to be **recursive** if it depends on itself.

View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:
 - repeat**
 - Find any view relation v_i in e_1
 - Replace the view relation v_i by the expression defining v_i
 - until** no more view relations are present in e_1

As long as the view definitions are not recursive, this loop will terminate

With Clause

- The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.
- **Find all accounts with the maximum balance**

```
with max_balance (value) as  
    select max (balance)  
    from account  
select account_number  
from account, max_balance  
where account.balance = max_balance.value
```

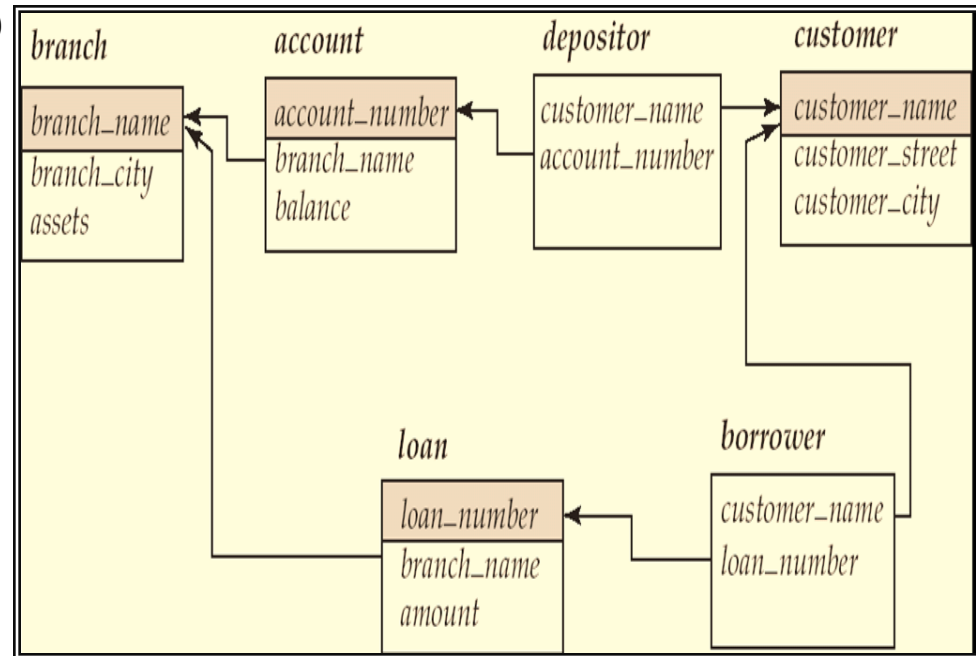
Complex Queries using With Clause

- Find all branches where the total account deposit is greater than the average of the total account deposits at all branches.

```
with branch_total (branch_name, value) as  
    select branch_name, sum (balance)  
    from account  
    group by branch_name
```

```
with branch_total_avg (value) as  
    select avg (value)  
    from branch_total
```

```
select branch_name  
    from branch_total, branch_total_avg  
    where branch_total.value >= branch_total_avg.value
```




Update of a View

- **Create a view of all loan data in the loan relation, hiding the amount attribute**

create view loan_branch **as**
select loan_number, branch_name
from loan

Add a new tuple to loan_branch

insert into loan_branch
values ('L-37', 'Perryridge')



This insertion must be represented by the insertion of the tuple

('L-37', 'Perryridge', null)

into the loan relation

Updates Through Views ..contd.

- **Some updates through views are impossible to translate into updates on the database relations**
 - **create view v as**
select loan_number, branch_name, amount
from loan
where branch_name = 'Perryridge'
insert into v values ('L-99','Downtown', '23')
- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation

Null Values

- It is possible for tuples to have a null value, denoted by null, for some of their attributes
- Null signifies an unknown value or that a value does not exist.
- The predicate **is null** can be used to check for null values.

Find all loan number which appear in the loan relation with null values for amount.

```
select loan_number  
from loan  
where amount is null
```

The result of any arithmetic expression involving null is null

Example: $5 + \text{null}$ returns null

However, aggregate functions simply ignore nulls

Null Values and Three Valued Logic

- **Any comparison with null returns unknown**
 - Example: $5 < \text{null}$ or $\text{null} <> \text{null}$ or $\text{null} = \text{null}$
- **Three-valued logic using the truth value unknown:**
 - OR: $(\text{unknown} \text{ or } \text{true}) = \text{true}$,
 $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$
 - AND: $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$,
 $(\text{false} \text{ and } \text{unknown}) = \text{false}$,
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$
 - NOT: $(\text{not unknown}) = \text{unknown}$
 - “P is unknown” evaluates to true if predicate P evaluates to unknown
- **Result of where clause predicate is treated as false if it evaluates to unknown**

Null Values and Aggregates

- Total all loan amounts

```
select sum (amount )  
from loan
```

- Above statement ignores null amounts
- Result is null if there is no non-null amount

- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes.

The where Clause . . . Contd.

- SQL includes a **between** comparison operator
- Find the loan number of those loans with loan amounts between Rs 90,000 and Rs 100,000 (that is, \geq Rs 90,000 and \leq Rs 100,000)

```
select loan_number  
  from loan  
 where amount between 90000 and 100000
```

